



MANGO
SOLUTIONS

Spark & R with sparklyr

Doug Ashton – Senior Data Scientist
Owen Jones – Data Scientist

#ODSC 



Agenda



- R and the Big Data Landscape
- The dplyr package for Spark
- Machine Learning with sparklyr



About Mango

www.mango-solutions.com
@MangoTheCat





MANGO
SOLUTIONS

R and the Big Data Landscape



R & Big Data



What is R?

- Fast: In memory calculations
- Easy:
 - Built for Data Science
 - Vibrant community building tools
 - Flexible, fast to develop
- But ...
 - Single Threaded
 - Poor scaling




The Yellow Elephant



Apache Hadoop

- HDFS (Hadoop Distributed File system)
 - Distributed, scalable file system
 - Store (and access) large files across machines
 - Written in Java
- MapReduce
 - Programming Model for processing parallelizable problems
 - Map > Shuffle > Reduce steps
 - Slow





Spark

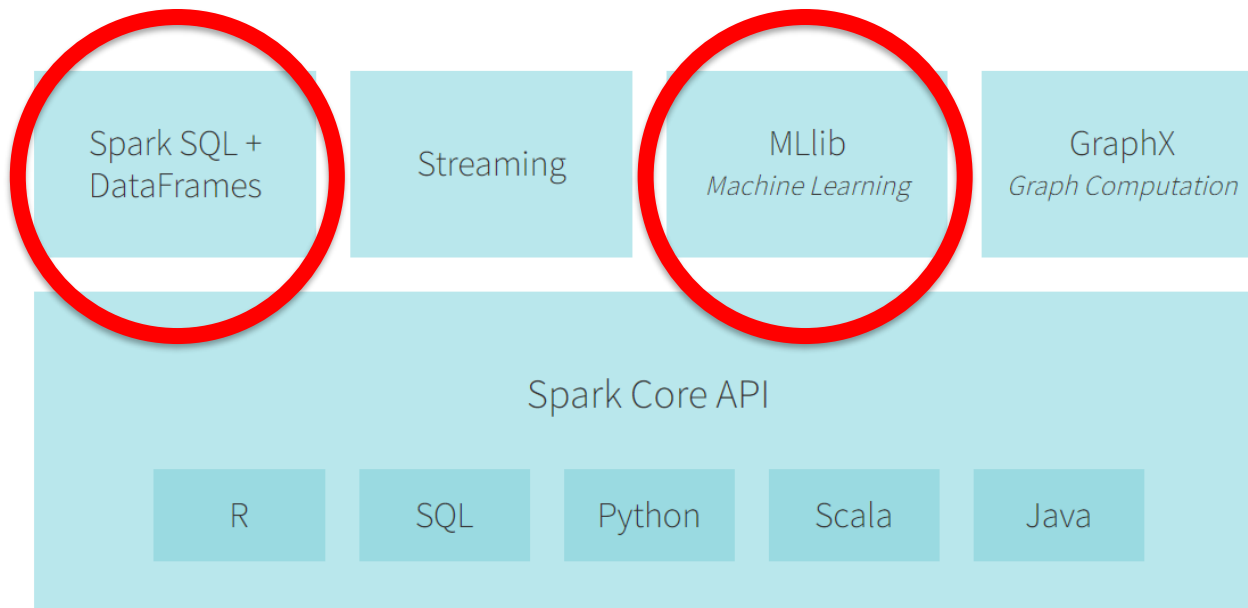


Apache Spark

- Fast: In-memory processing
- Easy: High level API designed for Data Science
- Scales: Distributed computation



Spark Ecosystem





Spark and R



Spark and R

- Originally supported Scala, Java and Python
- SparkR integrated into Spark as of v1.4
- Lets you create and work with MASSIVE data frames over a cluster of machines
- <https://spark.apache.org/docs/latest/sparkr.html>



SparkR

- Limited capability
- Unfamiliar code structures
- New implementations (e.g. of lm)



Sparklyr

- Rstudio with IDE integration
- Installer built in
- SparkSQL backend to **dplyr**
- Use Spark distributed ML library from R



What's in it for you?

- If you've ever run out of memory in R or need to interface with a Hadoop cluster, Spark and Sparklyr may be the solution for you.
- Allows you to create and work with vast datasets
- You can work interactively or create batch jobs, all using a familiar R/dplyr syntax



spark.mangodatalabs.com:8787

Apps ★ Bookmarks G + Pocket GBookmark Home Mango tools MIT - OCW | Comp... 7 Essential Lessons... GitHub · Where soft... Mango URLs Other Bookmarks

R File Edit Code View Plots Session Build Debug Profile Tools Help user75 Project: (None)

Go to file/function Adds

Console ~/

```
R version 3.3.3 (2017-03-06) -- "Another Canoe"
Copyright (C) 2017 The R Foundation for Statistical Computing
Platform: x86_64-redhat-linux-gnu (64-bit)

R is free software and comes with ABSOLUTELY NO WARRANTY.
You are welcome to redistribute it under certain conditions.
Type 'license()' or 'licence()' for distribution details.

Natural language support but running in an English locale

R is a collaborative project with many contributors.
Type 'contributors()' for more information and
'citation()' on how to cite R or R packages in publications.

Type 'demo()' for some demos, 'help()' for on-line help, or
'help.start()' for an HTML browser interface to help.
Type 'q()' to quit R.

> |
```

Environment History Spark

Import Dataset

Global Environment

Environment is empty

Files Plots Packages Help Viewer

New Folder Upload Delete Rename More

Home

	Name	Size	Modified
	R		



Creating a spark context

```
library(sparklyr)
```

```
config <- spark_config()
```

```
config$spark.ui.port <- "1234"
```

```
sc <- spark_connect(master = "spark://...:7077",  
                    config = config,  
                    app_name = "<your_name>")
```





MANGO
SOLUTIONS

The dplyr Package for Spark



Overview

- The Data
- Connect & Read
- The Core dplyr Functions
- SQL Functions, Joins and Sampling Data
- Save and Disconnect



Airlines Data

- Arrival and departure details for all commercial flights in US between October 1987 and April 2008.
- 120,000,000 records. **12 GB**
- stat-computing.org/dataexpo/2009/





MANGO
SOLUTIONS

Connect & Read



Connecting to Spark

```
sc <- spark_connect(  
    master = "...",  
    app_name = "my_name"  
)
```



Importing the Data

- Multiple options for connecting to data:
 - Copy data from R data frame
 - Import from file (CSV, JSON, PARQUET)
 - Using Spark SQL
 - From a Hive table



Importing the Data - Today

- We are going to read from a parquet file
 - Column store data format
- Can be read with

```
spark_read_parquet()
```



Read a Parquet File

```
airlines <-  
  spark_read_parquet(  
    sc,  
    name = "name_in_spark",  
    path = "path/to/file"  
  )
```



Importing the Data - Today

- We are also going to read from csv
- Can be read with

```
spark_read_csv()
```



The Data Object in R

- We don't work with data in R
- We work with a connection to the data
- dplyr will create Spark SQL to run
- The data will come to R (if we ask for it) after running the dplyr code



Exercise

- Ensure you have a connection to Spark
- Read the data and create a connection object to the table
- Read in the airports, carriers and plane-data CSV files
- What columns are in the airports data?
- Print the object in the console

(DO NOT `View` THE DATA YET)





MANGO
SOLUTIONS

Core dplyr Functions



The Core dplyr Functions

Function	Usage
<code>filter</code>	Filter the rows of a data set
<code>select</code>	Select columns from a data set
<code>mutate</code>	Add or manipulate columns in a data set
<code>arrange</code>	Sort the data
<code>summarise</code>	Generate summaries for columns in the data



A Quick Example

```
friday <- filter(flights, DayOfWeek == 5)
```

```
friday <- select(friday, -DayOfWeek)
```

```
friday <- mutate(friday,  
  Date = paste(Year, Month, DayofMonth,  
    sep = "-"))
```

```
select(friday, Year,  
  Month, DayofMonth, Date)
```



The Pipe Operator

- We can use dplyr as usual when working with spark
- This includes using the pipe operator (`%>%`) to simplify operations



Example with Pipes

```
flights %>%  
  filter(DayOfWeek == 5) %>%  
  select(-DayOfWeek) %>%  
  mutate(Date = paste(Year,  
    Month, DayofMonth, sep = "-"))  
  %>%  
  select(Year, Month,  
    DayofMonth, Date)
```



Exercise

- Create a query for the flights data where:
 - `depdelay > 15, depdelay < 240, dayofmonth == 15`
 - Columns: `year, month, arrdelay, depdelay, distance, uniquecarrier`
- Ensure your query is using the pipe operator



Lazy Execution

- The SQL query will be run at the last possible moment
- If you simply print results only 10 rows will be retrieved



How Do You Get All The Data?

```
flights %>%  
  filter(year > 2007) %>%  
  filter(depdelay == 240) %>%  
collect()
```



What is the Code Doing?

- The dplyr functions create a Spark SQL statement
- We can see the SQL statement by using

```
show_query()
```



Exercise

- Using the code you wrote in the last exercise
 - What does the Spark SQL query look like?
 - What are the dimensions of the data?





MANGO
SOLUTIONS

SQL Functions, Joins & Sampling



Translating to SQL

Operators	<code>+, -, *, /, %%, ^</code>
Mathematical functions	<code>abs, acos, cosh, sin, asinh, atan, atan2, atanh, ceiling, cos, cosh, cot, coth, exp, floor, log, log10, round, sign, sin, sinh, sqrt, tan, tanh</code>
Comparisons	<code><, <=, !=, >=, >, ==, %in%</code>
Booleans	<code>&, &&, , , !</code>
Aggregations	<code>mean, sum, min, max, sd, var, cor, cov, n</code>
Characters	<code>paste, tolower, toupper, nchar</code>
Casting	<code>as.double, as.integer, as.logical, as.character, as.date</code>



Joining and Sampling

Function	Usage
<code>group_by</code>	Apply a grouping to the data
<code>left_join,</code> <code>right_join,</code> <code>...</code>	Perform data joins (left, right, outer, ...)
<code>sample_n</code>	Sample n rows from the data



Joining

```
friday <- flights %>%  
  filter(DayOfWeek == 5) %>%  
  left_join(airports,  
    by = c(Dest = "iata")) %>%  
  filter(Origin %in%  
    c("SFO", "BOS"))
```



Sampling

```
friday %>%  
  sample_n(5)
```



Creating New Spark Data Frames

- When you join or sample you create a new data frame
- You may want to save this to use later
- We can do this with

```
sdf_register(r_obj, "spark_name")
```



Exercise

- Join the airport data to the flights data
- Filter the data to retain only flights originating in San Francisco and Boston
- Only retain data for Fridays
- Create a new Spark data frame containing this data





MANGO
SOLUTIONS

Save & Disconnect



Saving Your Results

- When we quit Spark our data won't be saved
- If we want to keep results we need to save them
- All of our import functions have equivalent write functions

```
spark_write_parquet()  
spark_write_csv()
```



Disconnecting (Don't Do This Now)

- Once we are done we need to disconnect from the spark instance

```
spark_disconnect(sc)
```





MANGO
SOLUTIONS

Machine Learning with sparklyr



Overview

- Mllib
- Data Preparation
- Model Training
- Model Evaluation





MANGO
SOLUTIONS

MLlib



MLlib Overview

“Apache Spark's scalable machine learning library”

- Data Preparation
 - Feature transformations, data partition
- Machine Learning Algorithms:
 - glm, Random Forest, GBM, K-means, Naive Bayes...
- Model Evaluation




MLlib in Context

- One of the four main libraries on top of core Spark
 - **Spark SQL**
 - Spark Streaming
 - **MLlib**
 - GraphX
- Shipped with Spark 0.8 (2013)



MLlib & Spark SQL

- Completely separate way to interact with DataFrames from Spark SQL.
- We can use them together!



Aside: Spark Core Structures

- RDD: Resilient Distributed Datasets
 - The original Spark data structure
 - Immutable general purpose structure
- Dataset:
 - New in Spark 1.6
 - More performant than RDDs
- DataFrame: A Dataset of rows
 - The only structure we'll be using



MLlib in sparklyr

- Three families of functions
- Feature Transformers: `ft_`
- Spark DataFrame Manipulation: `sdf_`
- Machine Learning: `ml_`





MANGO
SOLUTIONS

Data Preparation



Target Creation

- To demonstrate ft_ functions try the binarizer
- For classification can use threshold a continuous variable

```
iris_tbl <- copy_to(sc, iris)
```

```
ft_binarizer(iris_tbl,  
             input.col = "Sepal_Length",  
             output.col = "SL_Threshold",  
             threshold = 6)
```



Using MLlib in dplyr chain

- Where `mutate` maps to Hive (Spark SQL) functions, `sdf_mutate` uses MLlib
- Easier on the eye than raw `ft_` functions

```
flights %>%  
  mutate(ArrDelayDb = as.numeric(ArrDelay)) %>%  
  sdf_mutate(Late = ft_binarizer(ArrDelayDb,  
                                threshold = 30)) %>%  
  select(ArrDelay, Late)
```



Categorical (Factor) Data

- There is no “Factor” data type.
- Character data interpreted as factor
 - No need to convert for most ML algorithms
- Can converted using one hot encoding:
 - `ft_one_hot_encoder`
 - **Or** `ml_create_dummy_variables`

```
ml_create_dummy_variables(iris_tbl ,  
                           input = "Species")
```

Note: Need to drop reference var



Cut (bucketizer) 1

- `ft_bucketizer` works like R's `cut`

In R

```
iris %>%
```

```
  mutate(SL_Group = cut(Sepal.Length,  
                        breaks = 5))
```

In Spark

```
iris_tbl %>%
```

```
  sdf_mutate(SL_Group =  
    ft_bucketizer(Sepal_Length,  
                  splits = 4:8))
```



Cut (bucketizer) 2

In R

```
iris %>%  
  mutate(SL_Group = cut(Sepal.Length,  
                        breaks = 4:8))
```

In Spark

```
iris_tbl %>%  
  sdf_mutate(SL_Group =  
    ft_quantile_discretizer(Sepal.Length,  
                           n.levels = 5))
```



Other `ft_` functions

- `ft_one_hot_encoder`: Encoding categoricals
- `ft_tokenizer`: Splitting text into words
- `ft_string_to_index`: Category -> index and back again
- `ft_quantile_discretizer`: Faster, non-deterministic bucketizer
- `ft_vector_assembler`: Bring various vector cols back together



Exercise

Create a categorical feature in the friday dataset for morning and afternoon departures.

(hint `DepTime > 1200`)





MANGO
SOLUTIONS

Model Training



Available Algorithms

Function	Description
<u>ml_kmeans</u>	K-Means Clustering
<u>ml_linear_regression</u>	Linear Regression
<u>ml_logistic_regression</u>	Logistic Regression
<u>ml_survival_regression</u>	Survival Regression
<u>ml_generalized_linear_regression</u>	Generalized Linear Regression
<u>ml_decision_tree</u>	Decision Trees
<u>ml_random_forest</u>	Random Forests
<u>ml_gradient_boosted_trees</u>	Gradient-Boosted Trees
<u>ml_pca</u>	Principal Components Analysis
<u>ml_naive_bayes</u>	Naive-Bayes
<u>ml_multilayer_perceptron</u>	Multilayer Perceptron
<u>ml_lda</u>	Latent Dirichlet Allocation
<u>ml_one_vs_rest</u>	One vs Rest

Source: spark.rstudio.com

Final Data Prep

- Select **only** columns you want to use

```
select(friday , ArrDelay, DepTime)
```

- Last thing to do is partition into test and train

```
fl_part <- friday %>%
```

```
  select(ArrDelay, DepTime) %>%
```

```
  na.omit() %>%
```

```
  sdf_partition(train=0.8, test=0.2)
```

- Can have as many partitions as you like. Full cross validation not implement in R yet (cf H2O sparkling water).



Train the Model

- Similar to R. Can use formula interface.

```
model <- fl_part$train %>%  
  ml_linear_regression(ArrDelay ~ DepTime)
```

- Can't use interaction ($y \sim x1 : x2$) terms yet
- Alternatively, specify **Response** and **Features**

```
model <- fl_part$train %>%  
  ml_linear_regression(response = "ArrDelay",  
                        features = "DepTime")
```



Model Object

```
> class(model)
[1] "ml_model_linear_regression" "ml_model"
```

Works much like any other model

```
> summary(model)
```

Deviance Residuals::

Min	1Q	Median	3Q	Max
-81.260	-19.136	-8.682	5.231	1245.612

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)	
(Intercept)	-1.0703e+01	5.2566e-01	-20.361	< 2.2e-16	***
DepTime	1.5007e-02	3.7059e-04	40.495	< 2.2e-16	***



Scoring

- Use the `sdf_predict` function

```
scores <- sdf_predict(model,  
                      newdata = fl_part$test) %>%  
                      collect()
```



Elastic Net

- `ml_linear_regression` and `ml_logistic_regression` have lasso and ridge regression built in.
- Alpha: 0=ridge, 1=lasso
- Lambda: Sets strength of penalisation (use small numbers $\sim 10^{-3}$ (0 for normal lm))
- `ml_generalized_linear_regression` works like `glm`



Exercise

Fit a `ml_logistic` regression on the Friday dataset with a response of "Cancelled" predicted by departure time and destination state





MANGO
SOLUTIONS

Model Evaluation



Cross Validation

- MLlib does have cross validation routines but not yet available from sparklyr
- Can use H2O's cross validation with rsparkling



Area Under Curve

```
pred <- sdf_predict(model_class,  
                    fl_part$test)  
  
ml_binary_classification_eval(pred,  
    label = "Cancelled",  
    score = "probability",  
    metric = "areaUnderROC")
```





MANGO
SOLUTIONS

R on Spark



R on Spark Overview

- New in sparklyr 0.6
- Run arbitrary R code on all nodes
- Must have installed on all nodes
 - Packages are copied on first use (each session)
- Function with data frame in, data frame out



spark_apply

- A bit like lapply **over partitions**

```
iris2_tbl <- sdf_copy_to(  
  sc, iris, repartition = 2)
```

```
spark_apply(iris2_tbl,  
  function(e) {  
    data.frame(N=nrow(e))  
  },  
  names = "n",  
  packages = FALSE  
)
```



spark_apply

- Partitions are independent
- R packages are copied over once per session
 - This is super slow!
- If you add columns you have to declare them
- DataFrame -> function -> DataFrame



spark_apply

- Lets you use R only functions

```
spark_apply(  
  iris2_tbl,  
  function(e) {  
    dplyr::mutate(e,  
      Rpois = rpois(nrow(e), 3))  
  },  
  names = c(names(iris), "Rpois"),  
  packages = FALSE )
```





MANGO
SOLUTIONS

Concluding



Workshop Takeaways

- sparklyr evolving rapidly, things may change!
- Great tool for exploratory analysis on big datasets without needing to learn Scala/Java
- Built into Rstudio and Cloudera
- Can mix MLlib and Hive commands through dplyr
- Expect lots of new features soon

